

# HW5

Yigao Li

February 23, 2018

## 1. Feed Forward

(a)

(See Figure 1.)

(b)

$$h_1 = \tanh(1.2 + 4.2x_1 - 0.5x_2)$$

(c)

$$z = \tanh(5 - 8h_1 + 1.5h_2)$$

(d)

$$\begin{aligned} z &= \tanh(5 - 8h_1 + 1.5h_2) \\ &= \tanh(5 - 8\tanh(1.2 + 4.2x_1 - 0.5x_2) + 1.5\tanh(-30 + 20x_1 - 40x_2)) \end{aligned}$$

## 2. Multiple Solutions

Matrix of weights from input to hidden layer was  $B = \begin{pmatrix} 0 & 5 & 0 \\ 0 & 0 & 7 \end{pmatrix}$ ,  $w = \begin{pmatrix} -5 \\ 4 \\ 6 \end{pmatrix}$

Similar sets of weights lead to the same output:

1.  $B = \begin{pmatrix} 0 & -5 & 0 \\ 0 & 0 & 7 \end{pmatrix}$ ,  $w = \begin{pmatrix} -5 \\ -4 \\ 6 \end{pmatrix}$
2.  $B = \begin{pmatrix} 0 & 5 & 0 \\ 0 & 0 & -7 \end{pmatrix}$ ,  $w = \begin{pmatrix} -5 \\ 4 \\ -6 \end{pmatrix}$
3.  $B = \begin{pmatrix} 0 & -5 & 0 \\ 0 & 0 & -7 \end{pmatrix}$ ,  $w = \begin{pmatrix} -5 \\ -4 \\ -6 \end{pmatrix}$

Assume that  $g(y) = \tanh(y)$ ,  $g$  is an odd function. Output  $z = \tanh(-5 + 4\tanh 5x_1 + 6\tanh 5x_2)$ . Switching signs of coefficients of  $x_1$  and  $h_1$  does not change final output, similar to  $x_2$  and  $h_2$ :

$$\begin{aligned} z &= \tanh(-5 + 4\tanh 5x_1 + 6\tanh 5x_2) \\ &= \tanh(-5 - 4\tanh(-5x_1) - 6\tanh(-5x_2)) \\ &= \tanh(-5 - 4\tanh(-5x_1) + 6\tanh 5x_2) \\ &= \tanh(-5 + 4\tanh 5x_1 - 6\tanh(-5x_2)) \end{aligned}$$




Figure 1: Diagram of given neural network architecture

### 3. ANNs and Multiple Linear Regression

(a)

```
library(nnet)
```

```
## Warning: package 'nnet' was built under R version 3.4.3
```

```
set.seed(1234)
```

```
Advertising <- read.csv("Advertising.csv")
```

```
lin.model <- lm(Sales ~ TV + Radio + Newspaper, data = Advertising)
```

```
summary(lin.model)
```

```
##
```

```
## Call:
```

```
## lm(formula = Sales ~ TV + Radio + Newspaper, data = Advertising)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -8.8277 -0.8908  0.2418  1.1893  2.8292
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)  2.938889   0.311908   9.422  <2e-16 ***
```

```
## TV           0.045765   0.001395  32.809  <2e-16 ***
```

```
## Radio        0.188530   0.008611  21.893  <2e-16 ***
```

```
## Newspaper   -0.001037   0.005871  -0.177    0.86
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.686 on 196 degrees of freedom
## Multiple R-squared:  0.8972, Adjusted R-squared:  0.8956
## F-statistic: 570.3 on 3 and 196 DF,  p-value: < 2.2e-16

nnet.model <- nnet(Sales ~ TV + Radio + Newspaper, data = Advertising, size = 0, linout = TRUE,
                  skip = TRUE)

## # weights:  4
## initial value 102721.171716
## final value 556.825263
## converged

summary(nnet.model)

## a 3-0-1 network with 4 weights
## options were - skip-layer connections linear output units
## b->o i1->o i2->o i3->o
## 2.94 0.05 0.19 0.00

sum(residuals(lin.model)^2)

## [1] 556.8253

sum(residuals(nnet.model)^2)

## [1] 556.8253
```

Sum of squared errors are exactly the same for both methods.

(b)

```
Advert2 <- as.data.frame(scale(Advertising))
head(Advert2)

##      X      TV Radio Newspaper Sales
## 1 1 230.1  37.8      69.2  22.1
## 2 2  44.5  39.3      45.1  10.4
## 3 3  17.2  45.9      69.3   9.3
## 4 4 151.5  41.3      58.5  18.5
## 5 5 180.8  10.8      58.4  12.9
## 6 6   8.7  48.9      75.0   7.2

head(Advert2)

##           X           TV      Radio Newspaper      Sales
## 1 -1.719098  0.96742460  0.9790656  1.7744925  1.5481681
## 2 -1.701821 -1.19437904  1.0800974  0.6679027 -0.6943038
## 3 -1.684543 -1.51235985  1.5246374  1.7790842 -0.9051345
## 4 -1.667266  0.05191939  1.2148065  1.2831850  0.8581768
## 5 -1.649989  0.39319551 -0.8395070  1.2785934 -0.2151431
## 6 -1.632711 -1.61136487  1.7267010  2.0408088 -1.3076295

newAd <- Advert2
newAd$Sales <- Advertising$Sales
head(newAd)
```

```
##           X           TV           Radio Newspaper Sales
## 1 -1.719098  0.96742460  0.9790656  1.7744925  22.1
## 2 -1.701821 -1.19437904  1.0800974  0.6679027  10.4
## 3 -1.684543 -1.51235985  1.5246374  1.7790842   9.3
## 4 -1.667266  0.05191939  1.2148065  1.2831850  18.5
## 5 -1.649989  0.39319551 -0.8395070  1.2785934  12.9
## 6 -1.632711 -1.61136487  1.7267010  2.0408088   7.2
```

All data are converted to gaussian scaled values.

(c)

```
lin.model.scale <- lm(Sales ~ TV + Radio + Newspaper, data = newAd)
summary(lin.model.scale)
```

```
##
## Call:
## lm(formula = Sales ~ TV + Radio + Newspaper, data = newAd)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.8277 -0.8908  0.2418  1.1893  2.8292
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   14.0225     0.1192  117.655 <2e-16 ***
## TV              3.9291     0.1198   32.809 <2e-16 ***
## Radio          2.7991     0.1278   21.893 <2e-16 ***
## Newspaper     -0.0226     0.1279   -0.177    0.86
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.686 on 196 degrees of freedom
## Multiple R-squared:  0.8972, Adjusted R-squared:  0.8956
## F-statistic: 570.3 on 3 and 196 DF, p-value: < 2.2e-16
sum(residuals(lin.model.scale)^2)

## [1] 556.8253
```

## 4. MNIST Revisited

(a)

```
load("mnist_all.RData")
library(pROC)

## Warning: package 'pROC' was built under R version 3.4.3
## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
```

```

## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
index <- (train$y == 4 | train$y == 7)
df <- train$x[index,]
df.y <- train$y[index]
df <- as.data.frame(df)
df$y <- (df.y-4)/3 # 0 for number 4; 1 for number 7
a <- 711
b <- 298
var(df[,a])

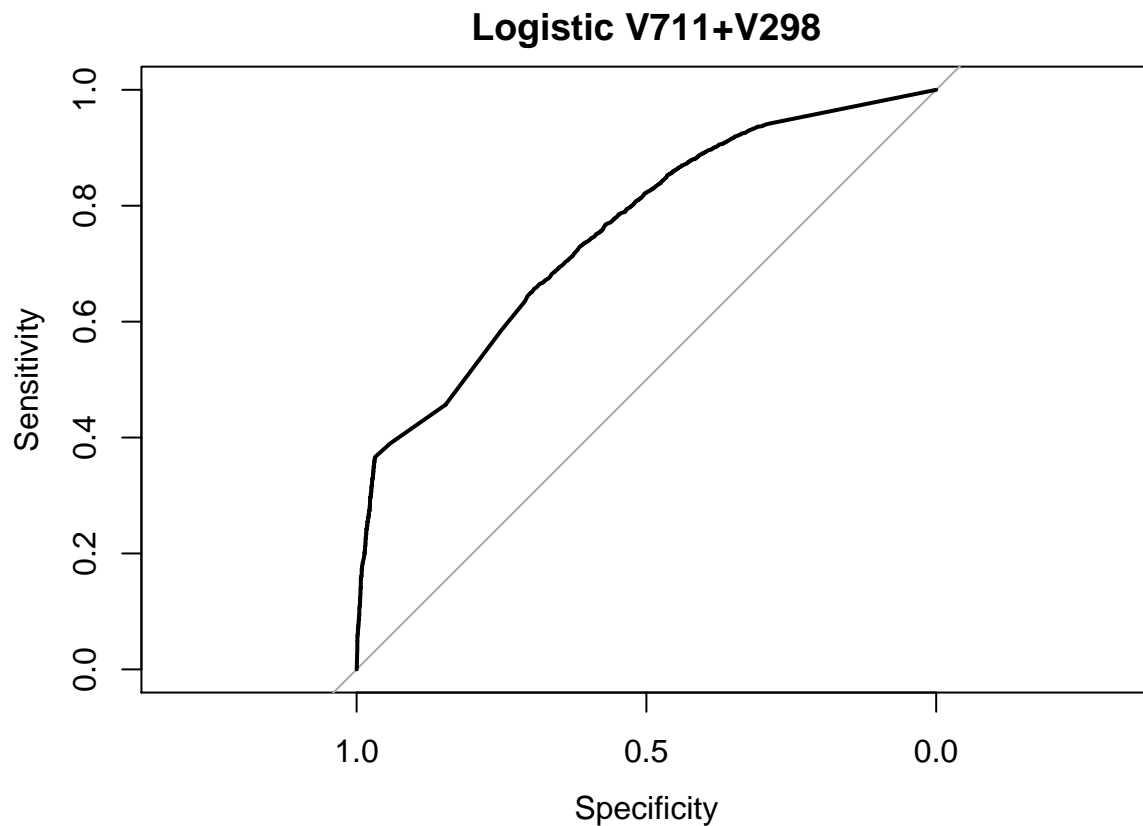
## [1] 6673.933
var(df[,b])

## [1] 11220.55
cor(df[,a], df[,b])

## [1] 0.009442966
model.1 <- glm(y ~ V711 + V298, data = df, family = binomial)
summary(model.1)

##
## Call:
## glm(formula = y ~ V711 + V298, family = binomial, data = df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.9213  -0.9918   0.1721   1.0906   1.6796
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.1308598  0.0383015  -29.52  <2e-16 ***
## V711         0.0158225  0.0004971   31.83  <2e-16 ***
## V298         0.0052907  0.0001985   26.66  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 16769  on 12106  degrees of freedom
## Residual deviance: 13856  on 12104  degrees of freedom
## AIC: 13862
##
## Number of Fisher Scoring iterations: 5
df$pred <- predict(model.1, type = "response")
myroc <- roc(df$y, df$pred)
plot(myroc, main = "Logistic V711+V298")

```



```
auc(df$y, df$pred)
```

```
## Area under the curve: 0.7543
```

(b)

```
set.seed(1)
model.2 <- nnet(y ~ V711 + V298, data = df, size = 1)
```

```
## # weights: 5
## initial value 2993.047137
## iter 10 value 2474.687781
## iter 20 value 2288.905982
## iter 30 value 2288.030887
## iter 40 value 2287.453102
## iter 50 value 2287.240971
## iter 60 value 2287.074669
## iter 70 value 2286.934932
## iter 80 value 2286.853979
## iter 90 value 2286.808223
## final value 2286.790916
## converged
```

```
summary(model.2)
```

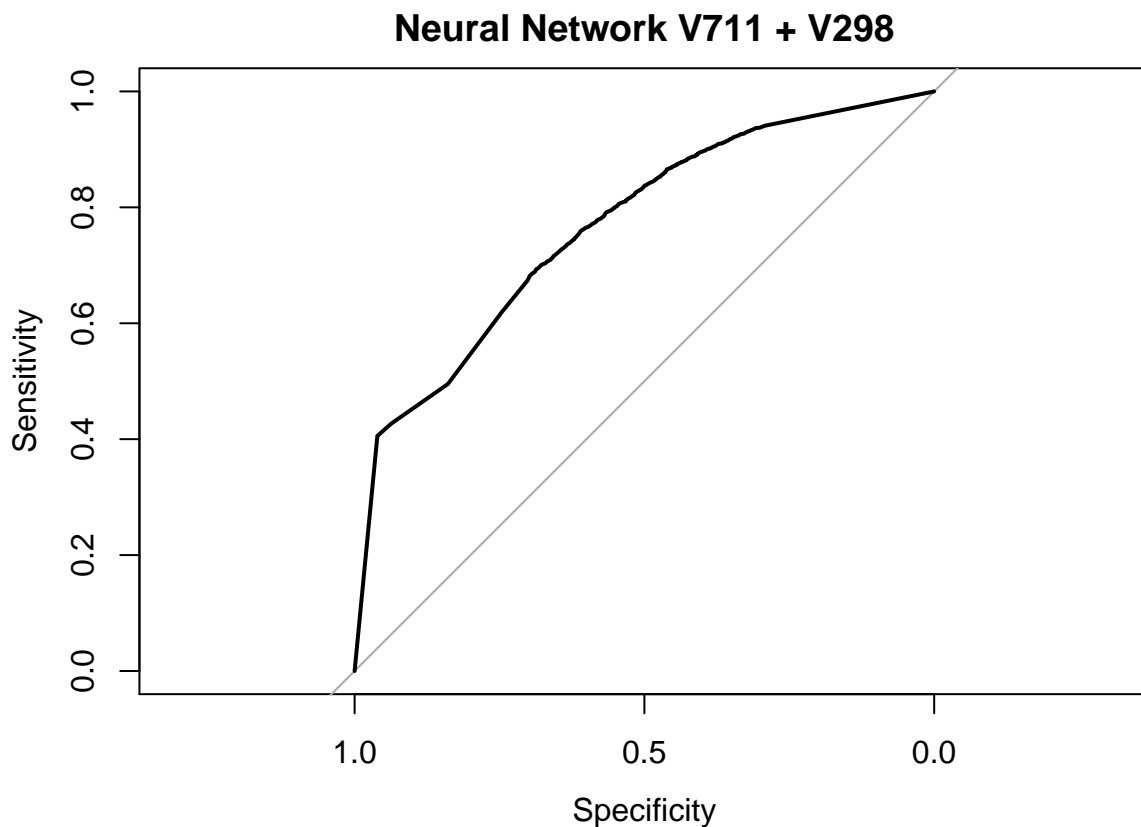
```
## a 2-1-1 network with 5 weights
```

```
## options were -
## b->h1 i1->h1 i2->h1
## -1.69 -8.02 0.00
## b->o h1->o
## 2.39 -23.65
```

```
df$pred <- predict(model.2, type = "raw")
myroc <- roc(df$y, df$pred)
```

```
## Warning in roc.default(df$y, df$pred): Deprecated use a matrix as
## predictor. Unexpected results may be produced, please pass a numeric
## vector.
```

```
plot(myroc, main = "Neural Network V711 + V298")
```



```
auc(df$y, df$pred)
```

```
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated
## use a matrix as predictor. Unexpected results may be produced, please pass
## a numeric vector.
```

```
## Area under the curve: 0.7663
```

Results from ANN slightly improves from logistic regression by 1.2% of training accuracy. 2 plots indicate that true positive rate is higher for ANN method with low false positive rate.

(c)

```
set.seed(2)

model.3 <- nnet(y ~ V711 + V298, data = df, size = 2)

## # weights:  9
## initial  value 2968.910960
## iter   10 value 2488.740396
## iter   20 value 2287.703570
## iter   30 value 2283.197239
## iter   40 value 2282.393015
## iter   50 value 2280.222002
## iter   60 value 2274.219847
## iter   70 value 2270.629453
## iter   80 value 2270.131971
## final   value 2269.942259
## converged

summary(model.3)

## a 2-2-1 network with 9 weights
## options were -
## b->h1 i1->h1 i2->h1
## -4.06  2.66  0.01
## b->h2 i1->h2 i2->h2
##  2.01  0.00  0.07
## b->o h1->o h2->o
## -8.93  3.28  8.31


df$pred <- predict(model.3, type = "raw")
myroc <- roc(df$y, df$pred)

## Warning in roc.default(df$y, df$pred): Deprecated use a matrix as
## predictor. Unexpected results may be produced, please pass a numeric
## vector.

plot(myroc, main = "Neural Network V711 + V298 (2 units in hidden layer)")
```



### Neural Network V711 + V298 (2 units in hidden layer)



```
auc(df$y, df$pred)
```

```
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated
## use a matrix as predictor. Unexpected results may be produced, please pass
## a numeric vector.
```

```
## Area under the curve: 0.7694
```

```
model.4 <- nnet(y ~ V711 + V298, data = df, size = 4)
```

```
## # weights: 17
## initial value 3060.513016
## iter 10 value 2373.051947
## iter 20 value 2306.585679
## iter 30 value 2278.954101
## iter 40 value 2276.525340
## iter 50 value 2274.768698
## iter 60 value 2270.848447
## iter 70 value 2270.235412
## iter 80 value 2269.882956
## iter 90 value 2269.521389
## iter 100 value 2269.080260
## final value 2269.080260
## stopped after 100 iterations
```

```
summary(model.4)
```


```
## a 2-4-1 network with 17 weights
```

```
## options were -
## b->h1 i1->h1 i2->h1
## -0.27  2.75 -1.51
## b->h2 i1->h2 i2->h2
##  5.24 -0.02 -0.02
## b->h3 i1->h3 i2->h3
## -2.90  2.67 -0.01
## b->h4 i1->h4 i2->h4
## -0.42 -0.15 -0.02
## b->o h1->o h2->o h3->o h4->o
##  1.47 -1.08 -1.90  2.90 -2.01
```

```
df$pred <- predict(model.4, type = "raw")
myroc <- roc(df$y, df$pred)
```

```
## Warning in roc.default(df$y, df$pred): Deprecated use a matrix as
## predictor. Unexpected results may be produced, please pass a numeric
## vector.
```

```
plot(myroc, main = "Neural Network V711 + V298 (4 units in hidden layer)")
```



```
auc(df$y, df$pred)
```

```
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated
## use a matrix as predictor. Unexpected results may be produced, please pass
## a numeric vector.
```

```
## Area under the curve: 0.7697
```

```
model.5 <- nnet(y ~ V711 + V298, data = df, size = 6)
```

```
## # weights: 25
## initial value 3728.023771
## iter 10 value 2337.455642
## iter 20 value 2276.452345
## iter 30 value 2269.746075
## iter 40 value 2268.582122
## iter 50 value 2267.717088
## iter 60 value 2267.244629
## iter 70 value 2266.949837
## iter 80 value 2266.799327
## iter 90 value 2266.750380
## iter 100 value 2266.609377
## final value 2266.609377
## stopped after 100 iterations
```

```
summary(model.5)
```


```
## a 2-6-1 network with 25 weights
## options were -
## b->h1 i1->h1 i2->h1
## -0.99 -0.05 1.41
## b->h2 i1->h2 i2->h2
## 0.75 -0.26 -0.23
## b->h3 i1->h3 i2->h3
## -6.19 2.21 5.23
## b->h4 i1->h4 i2->h4
## -10.52 6.24 0.03
## b->h5 i1->h5 i2->h5
## 7.86 -0.73 2.70
## b->h6 i1->h6 i2->h6
## -0.22 -0.17 -0.01
## b->o h1->o h2->o h3->o h4->o h5->o h6->o
## 0.12 2.55 -0.61 -1.65 2.85 -1.05 -1.95
```

```
df$pred <- predict(model.5, type = "raw")
myroc <- roc(df$y, df$pred)
```

```
## Warning in roc.default(df$y, df$pred): Deprecated use a matrix as
## predictor. Unexpected results may be produced, please pass a numeric
## vector.
```

```
plot(myroc, main = "Neural Network V711 + V298 (6 units in hidden layer)")
```

### Neural Network V711 + V298 (6 units in hidden layer)



```
auc(df$y, df$pred)
```

```
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated  
## use a matrix as predictor. Unexpected results may be produced, please pass  
## a numeric vector.
```

```
## Area under the curve: 0.769
```

The results did not improve. It is overfitting.